

(Q) Describe: ...

An Introduction to
Distributed Systems
SCC 311



(Q) Describe: Overview of the Session

1 Overview of the Session

What is a distributed system?
Examples of distributed systems
Why distributed systems?
Historical distributed systems
Recent trends in distributed
systems

(Q) Describe: What is a System?

2 What is a System?

A collection of hardware and software.

“A computer system consists of hardware and systems
software that work together to run application programs.”
Bryant and O’Hallaron

(Q) Describe: Tour of a Computer System

3 Tour of a Computer System

Registers, L1/L2/L3 caches, RAM, local disk, remote disk

A memory hierarchy

CPU running instructions

Carnegie Mellon

76

Example Memory

Hierarchy Regs

L1 cache

(SRAM)

Main memory

(DRAM)

Local secondary storage

(local disks)

Larger,

slower,

and

cheaper

(per byte)

storage

devices

Remote secondary storage

(e.g., Web servers)

Local disks hold files

retrieved from disks

on remote servers.

L2 cache

(SRAM)

L1 cache holds cache lines

retrieved from the L2 cache.

CPU registers hold words retrieved
from the L1 cache.

L2 cache holds cache lines

retrieved from L3 cache.

L0:

L1:

L2:

L3:

L4:

L5:

Smaller,
faster,
and
costlier
(per byte)
storage
devices

L3 cache
(SRAM)

L3 cache holds cache lines
retrieved from main memory.

L6:

Main memory holds disk
blocks retrieved from local
disks.

Carnegie Mellon

46

I/O Bus

Main

memory

I/O

bridgeBus interface

ALU

Register file

CPU chip

System bus Memory bus

Disk

controller

Graphics

adapter

USB

controller

Mouse Keyboard Monitor

Disk

I/O bus Expansion slots for

other devices such
as network adapters.



(Q) Describe: Tour of a Computer System

4 Tour of a Computer System

Registers, L1/L2/L3 caches, RAM, Disk
memory hierarchy
CPU with multiple cores
Operating system
Files, threads, ...
Assembly code
Compilers
Programs

(Q) Describe: What is a Distributed System?

5 What is a Distributed System?

“a collection of independent computers that appears to its users as a single coherent system”
Tanenbaum and van Steen

“one in which hardware or software components located at networked computers communicate and coordinate their actions only by passing messages”
Coulouris, Dollimore, Kindberg

“one that stops you getting work done when a machine you’ve never even heard of crashes”
Lamport

(Q) Describe: “a system designed to support the development of

6 “a system designed to support the development of

applications and services which can exploit a physical architecture consisting of multiple autonomous processing elements that do not share primary memory but cooperate by sending asynchronous messages over a communications network”
Blair and Stefani, 1998

What is a Distributed System?

(Q) Describe: In its broadest definition, a distributed system is one that

7 In its broadest definition, a distributed system is one that

comprises more than one computer with the goal of reaching a level of performance and/or providing a service that is quite difficult or infeasible to do on a single computer.

What is a Distributed System?



(Q) Describe: Examples of Distributed Systems

8 Examples of Distributed Systems

Financial trading

Support execution of trading transactions

Dissemination and processing of events

Web search

The web consists of over 6 billion web pages across 8.2 million servers with an average lifetime of about 2 months.

Modern web engines serve over 3.5 billion searches per day

>Major distributed systems challenges

Massively multiplayer online games

Online games (e.g. Fortnite, Among Us) support large numbers of users viewing and changing a common world

Need for very low latency coordination to support gameplay



(Q) Describe: Why Distributed Systems?

9 Why Distributed Systems?

Because the world is distributed

You want to book a hotel in Sydney, but you are in Lancaster

You want to be able to retrieve money from any ATM in the world,
but your bank is in London

An airplane has 1 cockpit, 2 wings, 4 engines, 10k sensors, etc.

Similarly railway networks, and other distributed transport systems

Because problems rarely hit two different places at the same
time

As a company having only one database server is a bad idea

Having two in the same room is better, but still risky

Because joining forces increases performance, availability, etc.

High Performance Computing, replicated web servers, etc.

(Q) Describe: Why Distributed Systems?

10 Why Distributed Systems?



DISTRIBUTED SYSTEMS

ACMQUEUE There's Just No Getting around It:
You're Building a Distributed System

Building a distributed system requires a methodical approach to requirements
Mark Cavage

Distributed systems are difficult to understand, design, build, and operate. They introduce exponentially more variables into a design than a single machine does, making the root cause of an application's problem much harder to discover. It could be said that if an application does not have meaningful SLAs (service-level agreements) and can tolerate extended downtime and/or performance degradation, then the barrier to entry is greatly reduced. Most modern applications, however, have an expectation of resiliency from their users, and SLAs are typically measured by "the number of nines" (e.g., 99.9 or 99.99 percent availability per month). Each additional 9 becomes harder and harder to achieve.

(Q) Describe: Is this important?

11 Is this important?



CLOUD

The Cloud team provides fundamental technology solutions around the automation and orchestration of continuous delivery pipelines targeting public cloud infrastructure.

MEDIA DISTRIBUTION

The Media Distribution team develops the software for geographically distributed caching servers that deal with the high throughput, high concurrency and low latency requirements of delivering BBC content to the audiences. If you've used iPlayer, there's a high likelihood that what you watched came from the in-house CDN that this team has developed.



Summary:

Facebook is seeking an experienced Software Engineer to join the Software Engineering (Infrastructure) team. The Software Engineering (Infrastructure) team builds large distributed components that run Facebook. Our code serves millions of requests per second and it does so with sub-second latency and in a fault tolerant manner. We handle everything from Facebook scale data storage, to synchronization and coordination of large server clusters, to providing a runtime environment for front end Facebook code. We are looking for candidates who share a passion for tackling complexity and building platforms that can scale through multiple orders of magnitude. This position is full-time and is based in our main office in New York, NY.



B B C



(Q) Describe: A Short History of DSs

12 A Short History of DSs

Distributed Computer Systems are largely concerned with:

Data processing/ management/ presentation (“computing” side)

Communication/ coordination (“distributed side”)

Those concerns existed well before computers were invented

Ancient civilisations needed efficient communication:

0 cf. the Postal Service of the Persian Empire (6th century BC),
the Roman roads (many still visible today), etc.

0 assuming a good infrastructure (roads, horses, staging posts)

0 max message speed: ~200 miles/day in the Persian system

Delays impose distributed organisations

0 Persian and Roman empires extended over 1000s of miles

Trust / secrecy / reliability issues

0 Am I sure Governor X is doing what he says he is?

This all has not really changed!

Things have only sped up.





DARIUS I THE GREAT

(Q) Describe: A Short History of DSs

13 A Short History of DSs

Computers are far more recent

The first “modern” computers appeared just after WWII

They were slow, bulky, and incredibly expensive

0 The ENIAC (1945), used by the US army: 30 tons, 170 m² footprint, 180 kilowatts, 18,000 vacuum tubes, and 5,000 additions/second (5KHz),

0 Price: \$500,000 (in US\$ of the time, would be >\$5,000,000 today)

And distributed computing is even more recent

For a long time, only very few computers were around anyway

No practical technology to connect them

This all changed in the 80's:

0 The rise of the micro-computers (PC, Mac, etc.)

0 The launch of the “Internet” (1982, TCP/IP), after 10 years of development

1940s 50s 60s 70s 80s 90s 00s 10s 20s

Fir

st

co

m

pu

ter

M

ain

fra

m

es

Mi

cro

-co

m

pu

ter

s

Th

e W

W

W

Gr

id
co
m
pu
tin
g
P2
P
co
m
pu
tin
g
Int
er
ne
t o
f T
hin
gs
Cl
ou
d c
om
pu
tin
g
Fo
g c
om
pu
tin
g
M
ini
-co
m
pu

ter

S



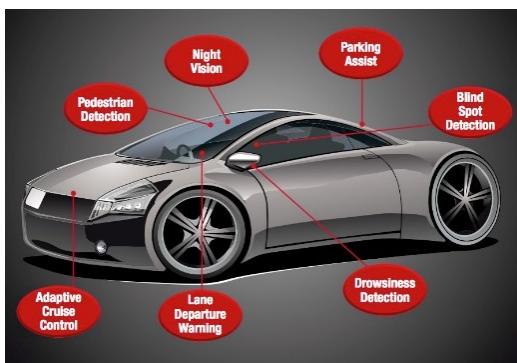
ARPANET GEOGRAPHIC MAP, OCTOBER 1969

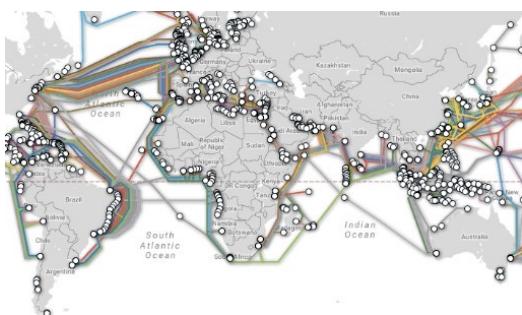


(Q) Describe: Recent Trends

14 Recent Trends

Pervasive networking and the modern Internet
Highly heterogeneous networking technologies
Fusion of services (telephony, broadcasting. . .)
Mobile and ubiquitous computing
Personal devices and embedded systems
Distributed multimedia systems
IPTV and Voice over IP services
Distributed computing as a utility
Grid, Cloud, and Fog computing
The pace of change is accelerating





(Q) Describe: Interactions

15 Interactions

What sort of interactions do we expect?

Web

Client: "I want this thing"

Server: "Here's the thing"

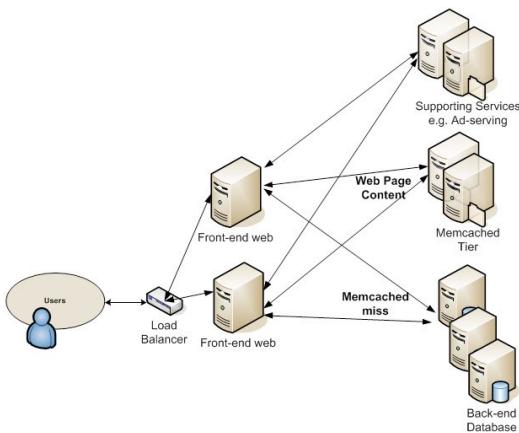
eBay

Alice: "I bid £10"

Bob: "I bid £12"

eBay: "Alice wins. Bob was late"

More on design patterns later (week 4).



(Q) Describe: Is this easy?

16 Is this easy?



(Q) Describe: But why is this so hard?

17 But why is this so hard?

A bank asks you to program their new ATM software

Central bank computer (server) stores account information

Remote ATMs authenticate customers and deliver money

A first version of the program

ATM: (ignoring authentication and security issues)

1. Ask customer how much money s/he wants
2. Send message with <customer ID, withdraw, amount >to bank server
3. Wait for bank server answer: <OK >or <refused >
4. If <OK >give money to customer, else display error message

Central Server:

1. Wait for messages from ATM: <customer ID, withdraw, amount >
2. If enough money withdraw money, send <OK >, else send <refused >

(Q) Describe: ATM Bank Server

18 ATM Bank Server

John: £500

But ...

What if the bank server crashes just after 2 and before 3?

What if the <OK> message gets lost? Takes too long to arrive?

What if the ATM crashes after 1, but before 4?

More on challenges in coming lectures.

<John, withdraw, £200 >

1

John: £300

-£200 2

<OK >

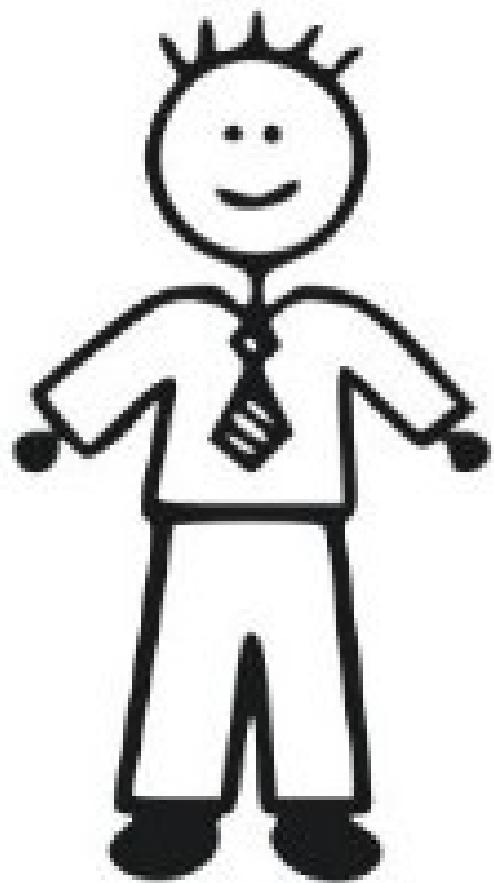
3

4

But why is this so hard?

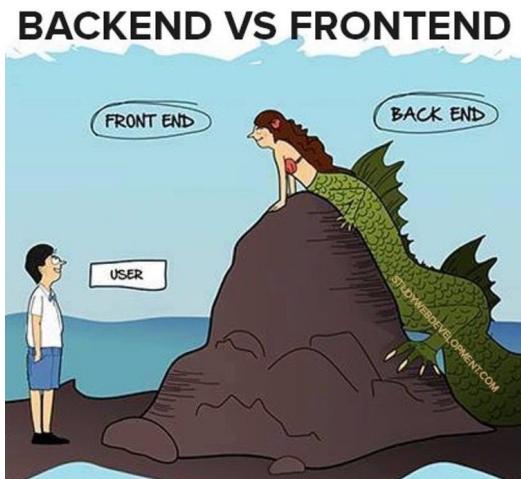






(Q) Describe: Is this easy?

19 Is this easy?



(Q) Describe: Expected Learning Outcomes

20 Expected Learning Outcomes

At the end of this session, you should be able to:

Define distributed computer systems

Explain why distribution is needed

Know a few everyday examples of distributed systems

Understand how they evolved into their current form

(Q) Describe: Additional Reading

21 Additional Reading

Required: CDKB, chapter 1 sections 1.1 – 1.4

also chapter 3 for revision

TvS, chapter 1

Computer Systems: A Programmer's Perspective, by Bryant
and O'Hallaron, chapter 1

Mark Cavage, "There's Just No Getting around It: You're
Building a Distributed System", ACM Queue, Vol. 11 No. 4,
April 2013. <http://queue.acm.org/detail.cfm?id=2482856>

(Q) Describe: ...

(Q) Describe: ...