

(Q) Describe: ...

Challenges of Distributed Systems
SCC311



(Q) Describe: Overview of the Session

1 Overview of the Session

Two Generals problem

Byzantine Generals problem

Fallacies of distributed computing

Middleware

(Q) Describe: Two Generals Problem

2 Two Generals Problem

Complication in distributed consensus

Two generals, each leading an army, want to capture a city. The attack is only successful if both armies attack together

Coordination is vital

Generals need to reach consensus in order to succeed in conquering the city.

Belisarius

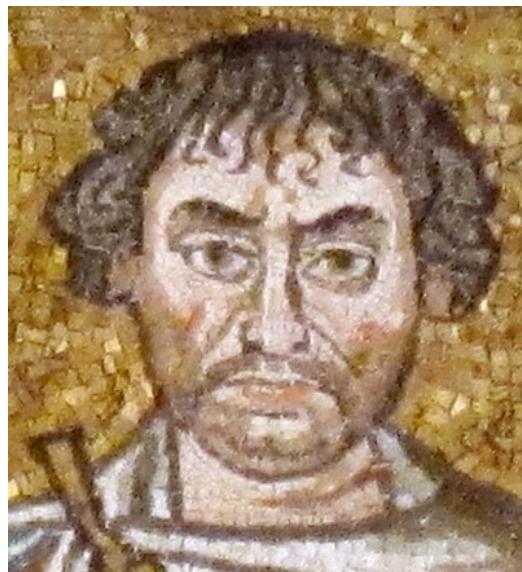
Narses

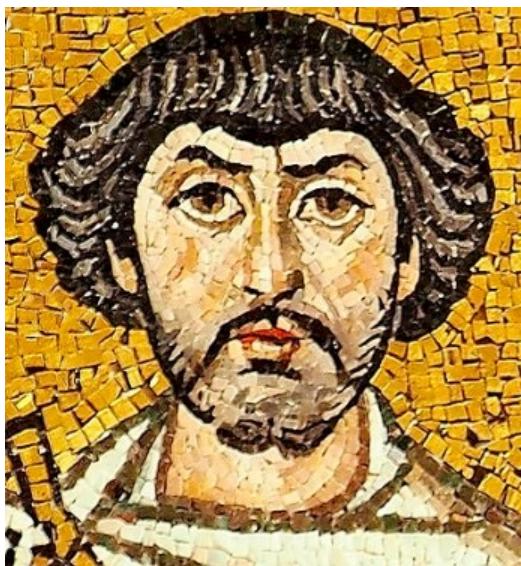
City

General1 General2

messengers

Attack or not Attack or not





(Q) Describe: Two Generals Problem

3 Two Generals Problem

Complication in distributed consensus

Two generals, each leading an army, want to capture a city. The attack is only successful if both armies attack together

Belisarius

Narses

City

General1 General2

messengers

Attack or not Attack or not

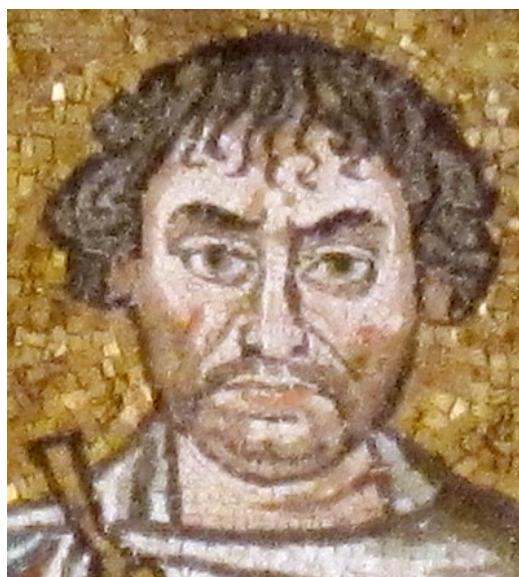
Army1 Army2 Result

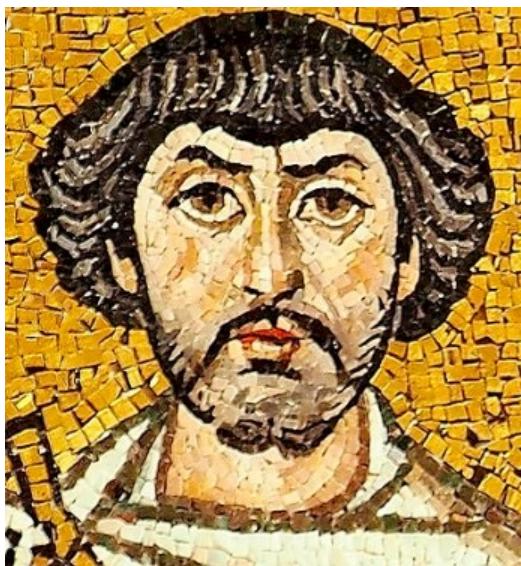
No Attack No attack Nothing happens

No Attack Attack Army2 defeated

Attack No attack Army1 defeated

Attack Attack City captured





(Q) Describe: Two Generals Problem

4 Two Generals Problem

Important challenges

Faulty network

Receipt of Acknowledgement is necessary

General1 General2

Attack on Oct 14 ?

Acknowledged!

General1 General2

Attack on Oct 14 ?

time time time time

(Q) Describe: How the generals should

5 How the generals should

decide?

Is it possible to come up with an algorithm that works all the time ?

Answer is no.

Possible approaches (none of them fully works):

Approach 1

General1 sends a lot of messages at once, and then attacks

Problem: if all the messengers are captured, army1 is defeated

Approach 2

General1 attacks only if an ACK is received from General 2

Now general2 will be captured if its response is intercepted or delayed

Problem: General2 has no idea if his ACK is received by General1

In the two generals problem, we assume a faulty network but processes (i.e., generals) are not faulty.

(Q) Describe: FLP Impossibility Result

6 FLP Impossibility Result

FLP impossibility result

Fischer, Lynch, and Patterson, Impossibility of Distributed Consensus with

One Faulty

Process, 1985

Asynchronous systems:

No bounds on message delivery time and process execution time

Clock drifts are also unbounded

When there is no upper bound on the time a process takes to complete its work and

respond, it's impossible to make the distinction between a process that is crashed and one

that is working (but taking very long to respond).

FLP shows that there is no guarantee for distributed processes to reach consensus in an

asynchronous environment, where it's possible for at least one process to crash.

Equivalently, it's not always possible to detect failure in an asynchronous system with

nodes crashing.

So let's consider the case where messages are delivered reliably (i.e., network is

not faulty) next.

(Q) Describe: Byzantine Generals Problem

7 Byzantine Generals Problem

We imagine that several divisions of the Byzantine army are camped outside an enemy city, each division has its own general. One commanding general (on behalf of the emperor) sends an order to the lieutenant generals. However, some of the generals, including the commanding general, may be traitors. The generals must have an algorithm to guarantee that:

- All loyal lieutenant generals decide upon the same plan of action.
- A small number of traitors cannot cause the loyal generals to adopt a bad plan.”
L. Lamport, R. Shostak, and M. Pease, “The Byzantine Generals Problem”, 1982

(Q) Describe: Byzantine Generals Problem

8 Byzantine Generals Problem

The Byzantine Generals Problem

LESLIE LAMPORT, ROBERT SHOSTAK, and MARSHALL PEASE
SRI International

Reliable computer systems must handle malfunctioning components that give conflicting information

to different parts of the system. This situation can be expressed abstractly in terms of a group of

generals of the Byzantine army camped with their troops around a nearby city. Communicating only

by messenger, the generals must agree upon a common battle plan. However, one or more of the m

ay be traitors who will try to confuse the others. The problem is to find an algorithm to ensure that

the loyal generals will reach agreement. It is shown that, using only oral messages, this problem is

solvable if and only if more than two-thirds of the generals are loyal; so a single traitor can confound

two loyal generals. With unforgeable written messages, the problem is solvable for any number of

generals and possible traitors. Applications of the solutions to reliable computer systems are then discussed.

Categories and Subject Descriptors : C.2.4. [Computer-Communication Networks]: Distributed

systems—network operating systems; D.4.4 [Operating Systems]: Communications Management—

network communication; D.4.5 [Operating Systems]: Reliability-fault tolerance

General Terms : Algorithms, Reliability

Additional Key Words and Phrases : Interactive consistency

/

1. INTRODUCTION

A reliable computer system must be able to cope with the failure of one or more

of its components. A failed component may exhibit a typical behavior which is

often overlooked -- namely, sending conflict in form
of a nation to different parts of
the system. The problem of coping with this type of failure
is explored abstractly
as the Byzantine Generals Problem. We derive them a
juxtaposition of the paper to a
discussion of this abstract problem and conclude by in-
dicating how our solutions
can be used in implement in a reliable computer system.
We argue that several ideas of the Byzantine Generals
problem can be used in implementing a reliable distributed system.
and general, The generals can
communicate with one another only by messenger. After
observing the enemy,
they must decide upon a common plan of action. However,
some of the generals
This research was supported in part by the National Aeronautics and
Space Administration under
contract NAS1-15428 Mod. 3, the Ballistic Missile Defense Systems
Command under contract
DASG60-78-C-0046, and the Army Research Office under contract DAAG29-
79-C-0102.
Authors' address: Computer Science Laboratory, SRI International,
333 Ravenswood Avenue, Menlo
Park, CA 94025.
Permission to copy without fee all or part of this material is granted
provided that the copies are not
made or distributed for direct commercial advantage, the ACM copyright
notice and the title of the
publication and its date appear, and notice is given that copying is by
permission of the Association
for Computing Machinery. To copy otherwise, or to republish, requires a fee
and/or specific
permission.
© 1982 ACM 0164-0925/82/0700-0382 \$00.75
ACM Transactions on Programming Languages and Systems, Vol. 4, No. 3,
July 1982, Pages 382-401.

The Byzantine Generals Problem

LESLIE LAMPORT, ROBERT SHOSTAK, and MARSHALL PEASE
SRI International

Reliable computer systems must be able to tolerate component failures by sending information to different parts of the system. This situation can be expressed abstractly in terms of a group of generals of the Byzantine army camped with their troops around an enemy city. Communicating only by messenger, the generals must agree upon a common battle plan. However, one or more of them may be traitors who will try to confuse the others. The problem is to find algorithms to ensure that the group reaches agreement even if some of the messages are forged or lost. The problem is solvable if and only if more than two-thirds of the generals are loyal; so a single traitor can confound two loyal generals. With unforgeable written messages, the problem is solvable for any number of generals and possible traitors. Applications of the solutions to reliable computer systems are then discussed.

Category and Subject Descriptors: C.0 [Computer Communication Networks]: Distributed Systems—network operating systems; D.4.1 [Operating Systems]: Communications Management—network communication; D.4.5 [Operating Systems]: Reliability—fault tolerance

General Terms: Algorithms, Reliability

Additional Key Words and Phrases: Interactive consistency

1. INTRODUCTION

A reliable computer system must be able to cope with the failure of one or more of its components. A failed component may exhibit a type of behavior that is often overlooked—namely, sending conflicting information to different parts of the system. The problem of coping with this type of failure is expressed abstractly as the Byzantine Generals Problem. We devote the major part of the paper to a discussion of this abstract problem and conclude by indicating how our solutions can be used in implementing a reliable computer system.

We imagine that several divisions of the Byzantine army are camped outside an enemy city, each division commanded by its own general. The generals can communicate with one another only by messenger. After observing the enemy, they must decide upon a common plan of action. However, some of the generals

This research was supported in part by the National Aeronautics and Space Administration under contract NASI-15428 Mod. 3, the Ballistic Missile Defense Systems Command under contract DASG001-80-C-0046, and the Army Research Office under contract DAAG29-79-C-0102.

Address reprint requests to Computer Science Laboratory, SRI International, 353 Ravenswood Avenue, Menlo Park, CA 94025.

Permission to copy without fee all or part of this material is granted provided that the copies are not made or distributed for direct commercial advantage, the ACM copyright notice and the title of the publication and its date appear, and notice is given that copying is by permission of the Association for Computing Machinery. To copy otherwise, or to republish, requires a fee and/or specific permission.

© 1982 ACM 0164-0925/82/0700-0082 \$00.75

ACM Transactions on Programming Languages and Systems, Vol. 4, No. 3, July 1982, Pages 382-401.

(Q) Describe: Byzantine Generals Problem

9 Byzantine Generals Problem

3 processes, commanding general is disloyal

Lieutenant1 Lieutenant2

Commander

??



(Q) Describe: Byzantine Generals Problem

10 Byzantine Generals Problem

3 processes, commanding general is disloyal

Lieutenant1 Lieutenant2

Commander

Retreat!Attack!



(Q) Describe: Byzantine Generals Problem

11 Byzantine Generals Problem

What should the loyal lieutenants do to follow a unified course of action?

Lieutenant1 Lieutenant2

Commander

Retreat!Attack!



(Q) Describe: Byzantine Generals Problem

12 Byzantine Generals Problem

What should lieutenants do ?

Lieutenants can communicate with one another and compare notes

Lieutenant1 Lieutenant2

Commander

Retreat!Attack!



(Q) Describe: Byzantine Generals Problem

13 Byzantine Generals Problem

Comparing notes (by sending messages to each other)

Remember this time network is not faulty

Lieutenant1 Lieutenant2

I was told

Attack

I was told

Retreat

Commander

Retreat!Attack!



(Q) Describe: Byzantine Generals Problem

14 Byzantine Generals Problem

Lieutenants can't agree on the commander's order so they do the default action which is retreat.

Commander

Lieutenant1 Lieutenant2

Retreat!Attack!

I was told

Attack

I was told

Retreat



(Q) Describe: Byzantine Generals Problem

15 Byzantine Generals Problem

Both scenarios are equivalent from Lieutenant1's point of view

Commander

Lieutenant1 Lieutenant2

Retreat!Attack!

I was

told

Attack

I was

told

Retreat

Commander

Lieutenant1 Lieutenant2

Attack!Attack!

I was

told

Attack

I was

told

Retreat





(Q) Describe: Byzantine Generals Problem

16 Byzantine Generals Problem

Both scenarios are equivalent from Lieutenant2's point of view

Commander

Lieutenant1 Lieutenant2

Retreat!Attack!

I was

told

Attack

I was

told

Retreat

Commander

Lieutenant1 Lieutenant2

Attack!Attack!

I was

told

Retreat

I was

told

Attack





(Q) Describe: Byzantine Generals Problem

17 Byzantine Generals Problem

With $d = 1$ traitor, the Byzantine Generals problem is not solvable with 3 generals

Consensus is not possible

You need at least 4 nodes! (see next slide)

Commander

Lieutenant1 Lieutenant2

Attack!Attack!

I was told

Attack

I was told

Retreat



(Q) Describe: Byzantine Generals Problem

18 Byzantine Generals Problem

With more than 3 participants, the strategy is still the same

All lieutenants compare notes with each other

Then follow the majority “vote”

If the majority was told “Attack”, then attack

If the majority was told “Retreat!”, then retreat

Commander

Lieutenant1 Lieutenant2

Attack!Attack!

I was told

Attack

I was told

Retreat



(Q) Describe: Byzantine Generals Problem

19 Byzantine Generals Problem

Communicating notes only work with one traitor

Commander

Lieutenant1

Retreat!Attack!

Lieutenant2 Lieutenant3 Lieutenant4 Lieutenant5 Lieutenant6 Lieutenant7



(Q) Describe: Byzantine Generals Problem

20 Byzantine Generals Problem

Let's say all lieutenants compare notes as before
Each lieutenant sends a message to every other lieutenant and indicates the order it received

Commander
Lieutenant1
Retreat!Attack!
I was told
Attack
Lieutenant2 Lieutenant3 Lieutenant4 Lieutenant5 Lieutenant6 Lieutenant7
I was told
Attack
I was told
Attack
I was told
Retreat
I was told
Retreat
I was told
Retreat





(Q) Describe: Byzantine Generals Problem

21 Byzantine Generals Problem

Let's say all lieutenants compare notes as before

Each lieutenant sends a message to every other lieutenant and indicates the order it received

Commander

Lieutenant1

Retreat!Attack!

I was told

Attack

Lieutenant2 Lieutenant3 Lieutenant4 Lieutenant5 Lieutenant6 Lieutenant7





(Q) Describe: Byzantine Generals Problem

22 Byzantine Generals Problem

How can the traitor lieutenant disrupt the majority vote and cause chaos ?

Commander

Lieutenant1

Retreat!Attack!

Lieutenant2 Lieutenant3 Lieutenant4 Lieutenant5 Lieutenant6 Lieutenant7

Attack:3

Retreat: 3





(Q) Describe: Byzantine Generals Problem

23 Byzantine Generals Problem

How can the traitor lieutenant disrupt the majority vote and cause chaos ?

Commander

Lieutenant1

Retreat!Attack!

Lieutenant2 Lieutenant3 Lieutenant4 Lieutenant5 Lieutenant6 Lieutenant7

I was told

Retreat

I was told

Attack





(Q) Describe: Byzantine Generals Problem

24 Byzantine Generals Problem

How can the traitor lieutenant disrupt the majority vote and cause chaos ?

Commander

Lieutenant1

Retreat!Attack!

Lieutenant2 Lieutenant3 Lieutenant4 Lieutenant5 Lieutenant6 Lieutenant7

Attack:4

Retreat: 3

Attack:4

Retreat: 3

Attack:4

Retreat: 3

Attack:3

Retreat: 4

Attack:3

Retreat: 4

Attack:3

Retreat: 4





(Q) Describe: Byzantine Generals Problem

25 Byzantine Generals Problem

The solution to this problem is more rounds of communication between lieutenants!

In addition to what they heard from the commander, lieutenants also share what they heard from each other.

Commander

Lieutenant1

Retreat!Attack!

Lieutenant2 Lieutenant3 Lieutenant4 Lieutenant5 Lieutenant6 Lieutenant7

L1 L2 L3 L4 L5 L6 L7

A A A A R R R

L1 L2 L3 L4 L5 L6 L7

R A A A R R R





(Q) Describe: Byzantine Generals Problem

26 Byzantine Generals Problem

For lieutenant2:

For lieutenant5:

Commander

Lieutenant1

Retreat!Attack!

Lieutenant2 Lieutenant3 Lieutenant4 Lieutenant5 Lieutenant6 Lieutenant7

L1 L3 L4 L5 L6 L7

A A A R R R

L1 L2 L3 L4 L6 L7

R A A A R R

I was told

Retreat

I was told

Attack





(Q) Describe: Byzantine Generals Problem

27 Byzantine Generals Problem

Then, the lieutenants exchange their vectors with each other
Each compose a table, where each row is a vector that is
exchanged in the previous round

Commander

Lieutenant1

Retreat!Attack!

Lieutenant2 Lieutenant3 Lieutenant4 Lieutenant5 Lieutenant6 Lieutenant7

L1 L2 L3 L4 L5 L6 L7

L1 X X X X X X

L2 A A A R R R

L3 A A A R R R

L4 A A A R R R

L5 R A A A R R

L6 R A A A R R

L7 R A A A R R

I was told

Retreat

I was told

Attack





(Q) Describe: Byzantine Generals Problem

28 Byzantine Generals Problem

Assumption: there is a default action in case loyal lieutenants can not agree on an action

Default = Retreat

Then, the lieutenants exchange their vectors with each other

Lieutenant5:

L1 L2 L3 L4 L5 L6 L7

L1 X X X X X X

L2 A A A R R R

L3 A A A R R R

L4 A A A R R R

L5 R A A A R R

L6 R A A A R R

L7 R A A A R R

? A A A R R R

Apply majority() to each column

(Q) Describe: Summary of Lamport's Algorithm

29 Summary of Lamport's Algorithm

For d traitors, there has to be more than $3d$ generals

For d faulty nodes, the algorithm requires a total of $d+1$ rounds of communication (d rounds of comparing notes + order from the commander)
Lamport's algorithm provides Byzantine Fault tolerance, given that you have more than $3d$ generals (processes)

Byzantine fault model: a process can behave arbitrarily (even maliciously as in the Byzantine Generals problem)

There are other fault models such as crash-stop (i.e., a process can stop responding), which are less challenging to deal with than the Byzantine fault model

We will cover fault tolerance in detail during week 5

Can we somehow make sure lieutenants don't forge orders that they didn't receive?

Public key cryptography can help with this – signed messages
However, traitors can collaborate and share their keys with each other
We will cover public key cryptography next week

(Q) Describe: Applications of Byzantine

30 Applications of Byzantine

Fault Tolerance

Space flight

Space shuttle had 4 computers for decision making

They used a majoritarian system in case of disagreements

Cryptocurrency: there are financial incentives to act maliciously

Problems stem from having multiple control systems

e.g. air traffic control, airline bookings, critical manufacturing processes, stock exchange, etc.

Lamport's algorithm is expensive! – $O(n^2)$ messaging overhead each round.

Alternative: accept uncertainty, just work around it

e.g. send a predetermined number of messages and assume at least one gets to the destination

Assume at most $d=1$ or 2 simultaneous bad actors

(Q) Describe: Further complications

31 Further complications

Heterogeneity
Access
Platform
Format
Administration
Changing nature
Increase/decrease in scale
Churn
Relocation
Failure



(Q) Describe: Fallacies of distributed computing

32 Fallacies of distributed computing

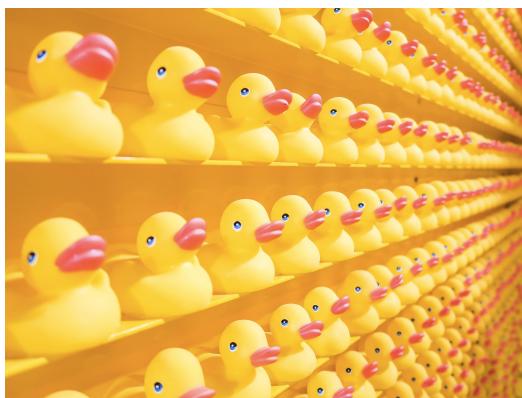
Asserted by Peter Deutsch (Sun Microsystems)

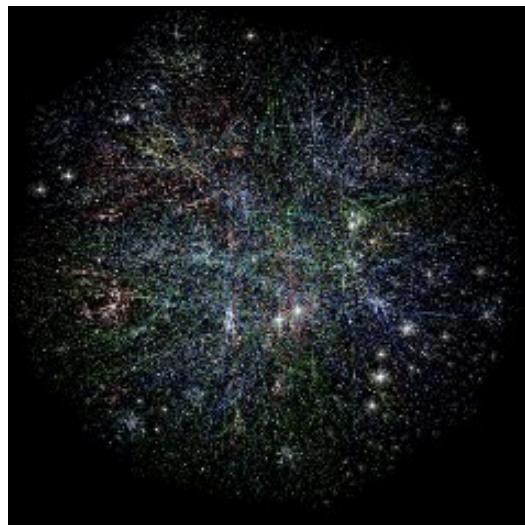
1. The network is reliable
2. Latency is zero
3. Bandwidth is infinite
4. The network is secure
5. There is one administrator
6. Transport cost is zero
7. The network is homogeneous
8. Topology doesn't change

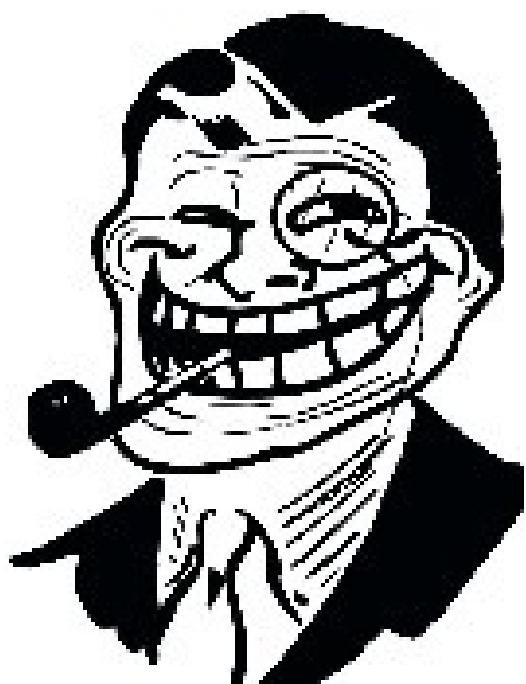












(Q) Describe: The Take-away Message?

33 The Take-away Message?

Challenging

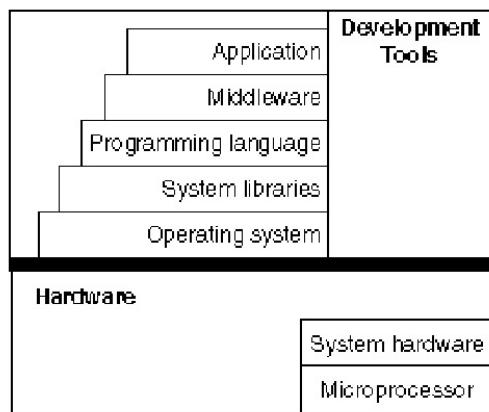
Developing good software is difficult

Developing good distributed software is even harder

To do this you need help!

Very hard to build such systems on bare-bones devices

Strong need for software platforms - >Middleware



(Q) Describe: Middleware

34 Middleware

Rationale

Provide a high level programming abstraction

Hide the complexity associated with distributed systems
(including the underlying forms of heterogeneity)

Machine A Machine B Machine C

Network

Network OS

Services

Middleware

Kernel KernelKernel

Network OS

Services

Network OS

Services

Distributed applications

uted Applic

(Q) Describe: Goals of a Middleware Platform

35 Goals of a Middleware Platform

Middleware provides protocols to support:

Resource sharing

The ability to access and share resources in a distributed environment

Distributed locking

Transparency

The ability to view a distributed system as if it were a single computer

Varying dimensions of transparency incl. location, access, migration, etc.

Openness

The offering of services according to standard rules (syntax and semantics)

Openness provides support for the key properties of:

◦ Portability: ability to run across different platforms

◦ Interoperability: having separate parts harmoniously work together

◦ Extensibility

The ability to be able to introduce new or modified functionality

(Q) Describe: Goals of a Middleware Platform

36 Goals of a Middleware Platform

Scalability: the ability to grow/shrink with respect to

Size

0 e.g. support massive growth in the number of users of a service

Geography

0 e.g. supporting systems across continents (dealing with latencies, etc.)

Administration

0 e.g. supporting systems spanning many different administrative organisations

Dependability/ Quality of Service

Security

0 Providing secure and authenticated channels, access control, key management, etc.

Fault tolerance

0 Providing highly available and resilient distributed applications and services

(Q) Describe: Styles of Middleware

37 Styles of Middleware

Client-server platforms

e.g. DCE (Distributed Computing Environment) – first!

Distributed object technology

e.g. CORBA, Java RMI

Component-based programming

e.g. Fractal, Enterprise Java Beans, OpenCOM

Microservice architecture

Loosely coupled, testable services using CI/CD

Others styles

Resource discovery platforms (e.g. Jini)

Group communication services (e.g. JGroups)

Publish-subscribe systems (e.g. JMS)

Distributed file systems, distributed transaction services, distributed document-based systems, agent-based systems, message-oriented middleware, P2P technologies, etc.

(Q) Describe: Expected Learning Outcomes

38 Expected Learning Outcomes

At the end of this session, you should:

Have an intuition of why realising good distributed systems is difficult

Byzantine generals problem as an example

Understand what the role of middleware in supporting the development of distributed applications and services

Goals of middleware systems

Some different styles of middleware

(Q) Describe: Additional Reading

39 Additional Reading

CDKB, chapter 1 section 1.5

also chapter 3 for revision

TvS, chapter 1, and sections 8.1-8.2

Lamport, Shostak, Pease, “The Byzantine Generals Problem”,

ACM Transactions on Programming Languages and Systems,

July 1982. <http://dl.acm.org/citation.cfm?id=357176>

Fischer, Lynch, Patterson, “Impossibility of Distributed Consensus with One Faulty Process”, Journal of ACM, April 1985.

<https://groups.csail.mit.edu/tds/papers/Lynch/jacm85.pdf>

Lamport, “Paxos made simple”,

<http://www.cs.utexas.edu/users/lorenzo/corsi/cs380d/past/03F/note s/paxos-simple.pdf>

Rotem-Gal-Oz, “Fallacies of Distributed Computing Explained”,

[\(2006\).](http://www.rgoarchitects.com/Files/fallacies.pdf)

(Q) Describe: ...

(Q) Describe: ...